

Coded Image Generation: [30pts]

The goal of this project is to obtain a simulated coded (or 'blurred') image from a given RGB image, Ground Truth Metric Depth Image and Point Spread Function of the phase mask.

Input RGB Image:

Input Metric Depth Image:

Output Coded RGB Image:

Mount Your Google Drive and Download PSF file from CodedVO Github:

```
# Step 1: Mount Google Drive to access and save files directly to your drive
# from google.colab import drive

# # Mount Google Drive
# drive.mount('/content/drive')

# Download and Save 'phasecam-psf-27.npy' from
https://github.com/naitri/CodedVO/blob/main/phasecam-psfs-27.npy

# Link to the datasets: (Use only RGB and Depth Images)
#
https://drive.google.com/drive/folders/12GrDxTBMaSlGeMRWycxmCQl01BHnC5-0
```

Import Environments:

```
import os           # Importing the os module for interacting with the
operating system
import cv2         # Importing OpenCV, a computer vision library
import numpy as np # Importing numpy for numerical operations
import torch       # Importing PyTorch for tensor operations and GPU
acceleration
import tqdm        # Importing tqdm for creating progress bars
import argparse    # Importing argparse for command-line argument
```

```

parsing
os.environ["OPENCV_IO_ENABLE_OPENEXR"] = "1"
import sys
# Enable support for reading and writing OpenEXR files in OpenCV

# Set the device to 'cuda' (GPU) if available, otherwise use 'cpu'
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(device)
# Enable cuDNN support for optimized GPU computation
torch.backends.cudnn.enabled = True

# Enable cuDNN benchmarking, which allows PyTorch to optimize
performance for the first time it's run
torch.backends.cudnn.benchmark = True

cuda

```

PyTorch Tensor to Numpy Conversion

```

def tensor_to_numpy_image(tensor: torch.Tensor) -> np.ndarray:
    """
    Convert a PyTorch tensor to a NumPy image.

    Args:
        tensor (torch.Tensor): Input tensor with shape (C, H, W),
                                where C is the number of channels,
                                H is the height, and W is the width.

    Returns:
        np.ndarray: Output image in NumPy format (H, W, C) with pixel
        values
                    in the range [0, 255].
    """

    # Move the channel axis from the first dimension (C, H, W) to the
    last dimension (H, W, C)
    out = tensor.moveaxis(0, -1)

    # Clamp the values to the range [0, 255], cast to uint8 type for
    image format,
    # move the tensor to CPU if it's on GPU, and convert to NumPy
    array.
    return torch.clamp(out, 0, 255).to(torch.uint8).cpu().numpy()

```

Get Depth Layers

Get depth layers from a depth map by quantizing depth values into discrete layers.

```
def get_depth_layers(depth_map: torch.Tensor, depth_layers:
    torch.Tensor) -> torch.Tensor:
    """
        Get depth layers from a depth map by quantizing depth values into
        discrete layers.

        Args:
            depth_map (torch.Tensor): Input depth map, a 2D or 3D tensor
            where each value represents the depth (distance) at a specific
            pixel.
            depth_layers (torch.Tensor): Defined depth layers, a 1D tensor
            that contains threshold values to separate the depth map into
            discrete layers.

        Returns:
            torch.Tensor: A tensor where each layer corresponds to a
            binary mask indicating whether each pixel belongs to a specific depth
            layer.
    """

    # Move the depth_layers tensor to the same device as depth_map for
    efficient computation.
    # This ensures that both tensors are compatible for GPU or CPU
    processing.
    depth_layers = depth_layers.to(depth_map.device)

    # torch.bucketize is a function that quantizes the depth_map
    values based on depth_layers.
    # Mathematically, it assigns each pixel in the depth_map to a
    "bin" (depth layer).
    # The input depth_map is a tensor of real values, and the
    depth_layers tensor contains discrete
    # boundary values (think of it as intervals that represent
    different ranges of depth).
    # For example:
    # If depth_layers = [0.5, 1.0, 1.5] and the depth_map contains
    values like [0.3, 0.6, 1.2, 1.8],
    # torch.bucketize assigns a bin index to each depth value:
    # 0.3 -> bin 0 (less than 0.5),
    # 0.6 -> bin 1 (between 0.5 and 1.0),
    # 1.2 -> bin 2 (between 1.0 and 1.5),
    # 1.8 -> bin 3 (greater than 1.5).
```

```
#####
#####
# WRITE A FUNCTION TO OBTAIN QUANTIZED DEPTH FROM 'depth_map' AND
'depth_layers' USING 'torch.bucketize' FUNCTION

#####
#####
quantized_depth = torch.bucketize(depth_map, depth_layers)
# print("Quantized Depth", np.shape(quantized_depth))

# torch.stack and list comprehension:
# For each depth layer, create a binary mask where the value is 1
if the pixel belongs to that layer
# and 0 otherwise. The quantized_depth == j comparison creates a
binary mask for each j in the
# range of the number of depth layers.

# Explanation of range(len(depth_layers)):
# len(depth_layers) returns the number of depth intervals.
# For each interval, we compare the quantized depth values with
the corresponding layer index (j).
# This results in a set of binary masks, each representing a
specific depth layer.

# Example:
# If quantized_depth = [0, 1, 2, 3] and len(depth_layers) = 4,
this will return four binary masks:
# Mask 0: True where quantized_depth == 0 (first depth layer),
False elsewhere.
# Mask 1: True where quantized_depth == 1 (second depth layer),
False elsewhere.
# Mask 2: True where quantized_depth == 2 (third depth layer),
False elsewhere.
# Mask 3: True where quantized_depth == 3 (fourth depth layer),
False elsewhere.

#####
#####
# WRITE A FUNCTION TO OBTAIN DEPTH LAYERS STACK FROM
'quantized_depth' AND 'depth_layers' USING 'torch.stack' FUNCTION

#####
#####
depth_layers_stack = torch.stack([(quantized_depth == i) for i in
range(len(depth_layers))], dim=0)
# print(depth_layers_stack)
```

```
return depth_layers_stack
```

Single PSF Convolution Operation:

Perform convolution of the input image with a single Point Spread Function (PSF).

PSF (Point Spread Function): This describes how a point source of light is blurred or spread by an imaging system. Each depth and channel can have a unique PSF that models this blur.

```
def single_psf_convolution(image: torch.Tensor, psfs: torch.Tensor,
    depth_idx: int, channel_idx: int, padding: int) -> torch.Tensor:
    """
        Convolve image with a single PSF.

    Args:
        image (torch.Tensor): Input image tensor, typically of shape
        (N, C, H, W) where:
            N - batch size, C - number of channels,
            H - height, W - width.
        psfs (torch.Tensor): Tensor of Point Spread Functions (PSFs),
        typically with shape (D, C, H, W) where:
            D - depth layers, C - number of channels,
            H - height, W - width.
        depth_idx (int): Index specifying which depth layer PSF to use
        for convolution.
        channel_idx (int): Index specifying which channel PSF to use
        for convolution.
        padding (int): Padding size to be applied during the
        convolution to control output size.

    Returns:
        torch.Tensor: The result of convolving the image with the
        selected PSF, output is of shape (N, C, H, W).
    """

    psfs = psfs.to(image.device).to(image.dtype) # Ensures psfs is on
    the GPU with the same dtype as the image
    # Write a function (Use torch's conv2d operator) that performs a
    2D convolution operation between the input image and the PSF.
    psf_conv = torch.nn.functional.Conv2d(image, psfs) # or we can use
    convolve_images()
    return psf_conv
```

Linear Convolution

The goal of this function is to apply different Point Spread Functions (PSFs) to the image based on the depth at each pixel. This allows for depth-dependent blur effects, where the blur at each pixel is determined by its corresponding depth value.

The linear model is given by: (For this model, the input RGBD image is quantized into K depth layers I_k , with $k=0$ being the furthest layer.)

$$b(\lambda) = \sum_{k=0}^{K-1} \text{PSF}_k(\lambda) * I_k(\lambda) + \eta$$

Follow: <https://www.computationalimaging.org/publications/deepopticsdfd/> [Equation 4] for notations.

```
def linear_convolution(image: torch.Tensor, depth_map: torch.Tensor,
psfs: torch.Tensor, depth_layers: torch.Tensor, padding: int) ->
torch.Tensor:
    """
    Perform depth-dependent convolution using a set of Point Spread
    Functions (PSFs), where each PSF is associated
    with a specific depth layer, and the final result is a depth-aware
    convolution of the image.

    Args:
        image (torch.Tensor): Input image tensor, typically of shape
        (C, H, W), where:
            C - number of channels (e.g., 3 for
            RGB), H - height, W - width.
        depth_map (torch.Tensor): Depth map tensor of shape (H, W),
        where each value represents the depth of a pixel.
        psfs (torch.Tensor): Point Spread Functions (PSFs) tensor of
        shape (D, C, H, W), where:
            D - number of depth layers, C - number of
            channels, H and W - height and width of each PSF.
        depth_layers (torch.Tensor): A tensor defining the discrete
        depth layers to segment the depth map.
        padding (int): Padding size for the convolution.

    Returns:
        torch.Tensor: The final convolved image after applying depth-
        aware convolutions.
    """

    # Move the psfs tensor to the same device as the input image for
    efficient computation,
    # ensuring that both tensors are compatible for GPU or CPU
    processing.
    psfs = psfs.to(image.device)
```

```

    # Generate binary masks (depth layers) from the depth map using
    the predefined depth_layers.
    # The result is a set of binary masks (depth_mask), each
    indicating where pixels belong to a specific depth layer.
    depth_mask = get_depth_layers(depth_map, depth_layers)
    # print("Depth Mask",np.shape(depth_mask))
    # Perform convolution for each channel (e.g., Red, Green, Blue in
    RGB images).
    # The final output will be a tensor of convolved images, stacked
    across channels.
    C,H,W = image.shape
    convolved_image = torch.zeros_like(image)

#####
#####
    # Write a function that performs a depth-dependent convolution on
    each color
    # channel of an image using point spread functions (PSFs)
    corresponding to
    # different depth layers, sums the results across depth layers,
    and then
    # stacks the convolved outputs for all channels into a final image
    tensor.
    # Hint: Use torch.stack, torch.sum and conv2d functions

#####
#####
    for i in range(len(depth_layers)):
        mask = depth_mask[i] # retrieving the current depth layer
        masked_image = image*mask.unsqueeze(0)
        # print("MaskedImage ",np.shape(masked_image))# according to
        pytorch site, not sure what it does
        convolved_channel = []
        for c in range(C):
            psf = psfs[i, c].unsqueeze(0).unsqueeze(0) # Shape (1, 1,
            H_psf, W_psf) for conv2d

            # Convolve masked image for each channel
            convolved = torch.nn.functional.conv2d(
                masked_image[c].unsqueeze(0).unsqueeze(0), # Shape (1, 1,
            H, W) for conv2d
                psf, # Shape (1, 1, H_psf, W_psf)
                padding=padding
            )

            # convolved_layer = torch.stack(convolved_channel,dim=0)
            # convolved_layer_1 = convolved_layer[1:len(convolved_layer)]
            # print(np.shape(convolved_layer_1))

```

```

    # print(np.shape(convolved_image))
    convolved_image[c] += convolved.squeeze(0).squeeze(0)

final_convolved_image = convolved_image
return final_convolved_image

```

Non Linear Convolution: (Optional)

A linear model can accurately reproduce defocus blur for image regions corresponding to a locally constant depth value. However, this approach is incapable of accurately modeling defocus blur at depth discontinuities.

Thus, we will adopt a nonlinear differentiable image formation model based on alpha compositing and combine it with our wavelength and depth-dependent PSF as:

$$b(\lambda) = \sum_{k=0}^{K-1} \tilde{l}_k \prod_{k'=k+1}^{K-1} (1 - \tilde{\alpha}_{k'}) + \eta$$

Follow: <https://www.computationalimaging.org/publications/deepopticsdfd/> [Equation 5] for notations.

```

def nonlinear_convolution():
    # Non linear code goes here...

    return final_convolved_image

```

Capture and Process all the image: (Nothing to write here)

Process all images and their corresponding depth maps in a folder using depth-dependent convolution.

```

def capture_image(img: np.ndarray, metric_depth: np.ndarray, psfs:
torch.Tensor, depth_layers: torch.Tensor, padding: int) -> np.ndarray:
    """
    Capture and process an image using depth-dependent convolution.

    Args:
        img (np.ndarray): Input RGB image as a NumPy array, with shape
        (H, W, 3).
        metric_depth (np.ndarray): Metric depth map as a NumPy array,
        with shape (H, W), where each value represents the depth at that
        pixel.
        psfs (torch.Tensor): Tensor of Point Spread Functions (PSFs)
        of shape (D, C, H, W), used for depth-dependent blurring.
        depth_layers (torch.Tensor): Tensor defining the depth layers
        that segment the depth map.
        padding (int): Padding size to use during the convolution

```


operation.

Returns:

np.ndarray: The final processed image after applying depth-dependent convolution, converted back to NumPy format.

```
"""  
  
# Convert the input RGB image from a NumPy array to a PyTorch  
tensor.
```

```
# `moveaxis(-1, 0)` changes the image from shape (H, W, 3) to (3,  
H, W) to match PyTorch's channel-first convention.
```

```
image = torch.from_numpy(img).moveaxis(-1,  
0).to(torch.float32).to(device)
```

```
# Convert the metric depth map from a NumPy array to a PyTorch  
tensor, and move it to the appropriate device (GPU or CPU).
```

```
depth = torch.from_numpy(metric_depth).to(device)
```

```
# Perform linear depth-dependent convolution on the input image  
using the given depth map, PSFs, and depth layers.
```

```
coded = linear_convolution(image, depth, psfs, depth_layers,  
padding)
```

```
# Convert the resulting convolved image (PyTorch tensor) back to a  
NumPy array in RGB format (H, W, 3).
```

```
return tensor_to_numpy_image(coded)
```

```
def process_folder(root: str, psfs: torch.Tensor, depth_layers:  
torch.Tensor, use_nonlinear: bool, is_blender: bool = False,  
scale_factor: float = 5000):  
    """
```

```
    Process all images and their corresponding depth maps in a folder  
    using depth-dependent convolution.
```

Args:

*root (str): Root directory containing "rgb" (image) and
"depth" folders.*

*psfs (torch.Tensor): Tensor of Point Spread Functions (PSFs)
to be applied for depth-dependent blurring.*

*depth_layers (torch.Tensor): Defined depth layers used for
segmenting the depth map.*

*use_nonlinear (bool): Flag indicating whether to apply non-
linear processing.*

*is_blender (bool): Flag indicating if depth files are in
Blender's EXR format.*

*scale_factor (float): Scale factor used for normalizing the
depth values.*

```
    """
```

```

    # Define paths for the depth and image folders within the root
    directory.
    depth_folder = os.path.join(root, "depth")
    image_folder = os.path.join(root, "rgb")
    os.environ["OPENCV_IO_ENABLE_OPENEXR"] = "1"

    # Create the output folder where processed images will be saved.
    output_folder = os.path.join(root, "Codedphasecam-27Linear-New")
    os.makedirs(output_folder, exist_ok=True) # Ensure the folder
    exists, create it if it doesn't.

    # Get the list of files in the image folder.
    files = os.listdir(image_folder)

    # Variable to keep track of the maximum depth value found across
    all depth maps.
    max_depth_value = 0

    # Padding size is half the width/height of the PSFs (assuming
    square filters).
    padding = psfs.shape[2] // 2

    # Loop through each file in the image folder, displaying progress
    using tqdm.
    for idx, file in tqdm.tqdm(enumerate(files), total=len(files),
    desc=root):

        # Construct the full paths for the image and corresponding
        depth files.
        image_file = os.path.join(image_folder, file)
        depth_file = os.path.join(depth_folder, file).replace(".png",
        ".exr") if is_blender else os.path.join(depth_folder, file)

        # Read the image file (in BGR format since OpenCV loads in BGR
        by default).
        image_bgr = cv2.imread(image_file)

        # Convert the image from BGR to RGB format (as needed for
        processing).
        image = cv2.cvtColor(image_bgr, cv2.COLOR_BGR2RGB)

        # Read the depth file, handling both EXR (Blender) and other
        formats.
        raw_depth = cv2.imread(depth_file, cv2.IMREAD_ANYCOLOR |
        cv2.IMREAD_ANYDEPTH)

        # Check if the depth file is missing or unreadable, skip
        processing for that file if so.
        if raw_depth is None:
            print(f"{file} is missing a depth file")

```

```

        continue

        # If using Blender's EXR format, extract the depth from the
        first channel.
        # Otherwise, normalize the depth using the provided scale
        factor.
        metric_depth = raw_depth[:, :, 0] if is_blender else raw_depth
        / scale_factor

        # Apply depth-dependent convolution to the image, either
        linear or non-linear depending on the flag.
        coded_image_rgb = capture_image(image.astype(np.float32),
        metric_depth, psfs, depth_layers, padding)

        # Convert the processed image back to BGR format for saving
        (as OpenCV saves in BGR by default).
        coded_image_bgr = cv2.cvtColor(coded_image_rgb,
        cv2.COLOR_RGB2BGR)

        # Save the processed image to the output folder.
        cv2.imwrite(os.path.join(output_folder, file),
        coded_image_bgr)

        # Update the maximum depth value encountered in this folder.
        max_depth_value = max(max_depth_value, np.max(metric_depth))

        # Print the maximum depth value found in the folder for reference.
        print(f"Max Depth Value in the folder: {max_depth_value}")

```

Main Function:

Put everything together! Modify the Image and PSF file path according to your needs.

Note that scale_factor relies on the camera hardware and is different for different datasets.

Note: Our Point Spread Functions (PSFs) correspond to discretized depth layers using a 23×23 Zernike parameterized phase mask, with the depth range discretized into 27 bins within the interval of [0.5, 6] meters, with a focal distance of 85 cm.

```

def main():
    # Define and configure an argument parser (currently commented
    out).
    # The argument parser allows users to specify the root directory,
    whether the depth images are in Blender's EXR format,
    # and the scale factor for depth normalization through command-
    line arguments.
    parser = argparse.ArgumentParser(description="Process images with
    depth-dependent processing.")

```

```

    parser.add_argument("--root", type=str, required=True, help="Root
directory containing the datasets.")
    parser.add_argument("--is_blender", action="store_true", help="Use
Blender's EXR depth format.")
    parser.add_argument("--scale_factor", type=float, default=5000,
help="Scale factor for depth normalization.")
    args = parser.parse_args(["--root", "C:/Adv Computer
Vision/P2/UMD-CodedV0-dataset/DiningRoom/", "--is_blender", "--
scale_factor", "1"])
    #Following arguments are for running Livingroom and NYU Data
    # args = parser.parse_args(["--root", "C:/Adv Computer
Vision/P2/train/LivingRoom1/", "--is_blender", "--scale_factor", "1"])
    # args = parser.parse_args(["--root", "C:/Adv Computer
Vision/P2/train/nyu_data/nyu_data/", "--scale_factor", "1000"])

    # Print the device being used (e.g., GPU or CPU), which was set
earlier in the code using `torch.device()`.
    print(device)
    # Define the root directory where the image and depth datasets are
located.
    # This is hardcoded in the script but could be passed as an
argument using the argument parser (currently commented).
    # root = '/content/drive/MyDrive/train/LivingRoom1'

    root = args.root # Root directory from arguments
    is_blender = args.is_blender # Whether to use Blender's EXR depth
format
    scale_factor = args.scale_factor
    # root = 'C:/Adv Computer
Vision/P2/UMD-CodedV0-dataset/DiningRoom/'

    # Flag indicating whether the depth maps are in Blender's EXR
format.
    # is_blender = True

    # Define the scale factor for depth normalization (used when the
depth maps are not in EXR format).
    # scale_factor = 5000 # This scale is different for different
datasets
    # It is 1000 for NYUv2 dataset
    # It is 5000 for ICL-NUIM dataset
    # It is 1 for all other dataset provided (UMD-CodedV0/DiningRoom
etc. datasets)
    # Depth sensors and datasets typically store depth as integer or
floating-point
    # values. However, these values are often not in direct metric
units (like meters)
    # because it would require storing large floating-point values,
    # which increases storage size and precision issues.

```

```

# Load the Point Spread Functions (PSFs) from a `.numpy` file.
# psf_path = '/content/drive/MyDrive/phasecam-psfs-27.npy'
psf_path = './phasecam-psfs-27.npy'

# Define the depth layers, which divide the depth map into
segments based on depth values.
# `torch.linspace` creates a tensor with 27 evenly spaced depth
layers between 0.5 and 6 units of depth.

#####
#####
# TODO: depth_layers = ...

#####
#####
depth_layers = torch.linspace(0.5, 6.0, 27).to(device)
# Load the PSFs from the `.numpy` file and adjust their axis order
using `np.moveaxis`.
# This is done to match the format expected by the convolution
function.
psfs = torch.from_numpy(np.moveaxis(np.load(psf_path), -1, 1))

# Call the `process_folder` function to process all images in the
folder.
# The parameters passed include the root directory, loaded PSFs,
defined depth layers,
# a flag for whether to use non-linear processing, the Blender
depth format flag, and the scale factor.
process_folder(
    root=root,
    psfs=psfs,
    depth_layers=depth_layers,
    is_blender=is_blender,
    use_nonlinear = False, # Use Blender's EXR depth format if set
to True.
    scale_factor=scale_factor, # The scale factor for depth
normalization.
)

# Run the main function when the script is executed.
main()

cuda

C:/Adv Computer Vision/P2/UMD-CodedV0-dataset/DiningRoom/: 100%|
██████████| 1000/1000 [04:31<00:00, 3.69it/s]

Max Depth Value in the folder: 65504.0

```

Submission Guidelines:

You are required to submit:

1. The compiled version of this colab file with results
2. PDF Report (Combined with Part 2)
3. Final Coded Image Output using Linear Convolution on the following image set (Images: '1.png', '10.png', '20.png', '30.png', '40.png')

To confirm your results are correct, compare the results of your linear model with the following images:

<https://drive.google.com/drive/folders/1m0ihEGnDnOtedTpl8VRJ8miD25Bxqxa-?usp=sharing>

(Note that: Do not compare with Non-Linear nyu_data output given here:

<https://drive.google.com/drive/folders/12GrDxTBMaSlGeMRWycxmCQl01BHnC5-O>)

Note: Except the pixels in the final coded image that are far away to be black something like this:

https://drive.google.com/file/d/14Kg2BCs-Mie5Y-udDanNJPaz0PGae_nE/view?usp=sharing